LIGHTWEIGHT MULTIMODAL NEURAL NETWORK USING MOBILE DIAGNOSTICS

A PROJECT REPORT

Submitted by

PRERNA GUPTA (24BDA70020) PRACHI (24BDA70046)
YASHPREET KOUR (24BDA70043) UCHIT (24BDA70054)
PARTH SHARMA (24BDA70014)

in partial fulfillment for the award of the degree of

BACHELORS OF ENGINEERING

IN

CSE-DATA SCIENCE

Chandigarh University

November,2025

BONAFIDE CERTIFICATE

Certified that this project report **"LIGHTWEIGHT MULTIMODAL NEURAL NETWORKS USING MOBILE DIAGNOSTICS"** is the bonafide work of "**PRERNA GUPTA (24BDA70020), PRACHI (24BDA70046) ,YASHPREET KOUR (24BDA70043) , UCHIT (24BDA70054), PARTH SHARMA (24BDA70014)"**

who carried out the project work under my/our supervision.

SIGNATURE

Dr. Aman Kaushik

HEAD OF THE DEPARTMENT

AIT-CSE

SIGNATURE

Ms. Somdatta Patra

SUPERVISOR

Assistant Professor

AIT-CSE

Submitted for the project viva-voce examination held on

INTERNAL EXAMINER **EXTERNAL EXAMINER**

TABLE OF CONTENTS

ABSTRACT

The increasing demand for real-time, accessible, and intelligent healthcare solutions has driven the need for lightweight artificial intelligence (AI) systems capable of running efficiently on mobile and wearable devices. This project presents a **Lightweight Multimodal Neural Network (LMNN)** designed for **mobile health diagnostics**, integrating multiple input modalities such as **medical images, audio signals, and gesture data** to enhance diagnostic accuracy and robustness.

The proposed framework combines **Convolutional Neural Networks (CNNs)**, **Gated Recurrent Units (GRUs)**, and **attention-based fusion mechanisms** to efficiently process heterogeneous data while minimizing computational overhead. To ensure compatibility with limited-resource environments, **model compression techniques** such as **pruning** and **quantization** were applied, reducing model size to **approximately 12 MB** and achieving inference latency of **under 100 milliseconds** on standard mobile processors.

Experimental evaluation demonstrated that the LMNN achieved **over 92% diagnostic accuracy**, outperforming unimodal approaches while maintaining low energy consumption and high reliability, even in the presence of missing or noisy modalities. The system's design enables **on-device, privacy-preserving diagnostics**, making it suitable for applications in **telemedicine, wearable health tracking, and emergency care**.

This research contributes a scalable, energy-efficient, and privacy-aware AI framework that bridges the gap between advanced multimodal diagnostics and mobile deployment, advancing the goal of **personalized, real-time healthcare for all**.

CHAPTER 1.
INTRODUCTION

1.1. Client Identification/Need Identification/Identification of relevant Contemporary issue

Issue Identification:
With the exponential growth of mobile health (mHealth) technologies, there is a rising demand for AI-driven diagnostic systems capable of operating in real-time on resource-limited mobile and wearable devices. Traditional multimodal neural networks are highly resource-intensive, making them unsuitable for such environments.

Justification through Statistics and Documentation:

- According to the *World Health Organization (2024)*, over 60% of healthcare applications rely on mobile platforms for remote diagnostics and monitoring.
- *IEEE IoT Journal (2024)* reports that lightweight AI models can reduce computational cost by up to 45%, enabling real-time inference under 100 ms on mobile devices.
- *National Library of Medicine (2024)* datasets such as MedPix 2.0 and MHAD highlight the global shift toward multimodal diagnostic systems integrating image, audio, and gesture data.

Problem Need (Consultancy Aspect):
Hospitals, healthcare startups, and wearable technology companies need energy-efficient and accurate AI models for mobile diagnostic applications that can operate offline, on-device, and in real time — improving accessibility for remote or under-resourced regions.

Need Justified through Reports/Surveys:
Surveys from *IEEE Access (2023)* and *ACM Transactions on Computing for Healthcare (2025)* indicate that the lack of optimized AI architectures for mobile environments is one of the top technical bottlenecks in deploying AI healthcare tools.

Thus, the need for a unified, lightweight multimodal neural network framework for mobile diagnostics is a validated and contemporary research problem.

1.2. Identification of Problem

Broad Problem Statement:
Modern multimodal AI models for medical diagnostics are computationally expensive, energy-inefficient, and unsuitable for real-time execution on mobile devices.
The challenge is to design an optimized multimodal neural network that maintains diagnostic accuracy while minimizing power, memory, and latency costs for deployment on mobile and wearable healthcare systems.

1.3. Identification of Tasks

To systematically address the problem, the project can be divided into the following tasks:

Task 1: Problem Identification and Literature Review

- Study existing multimodal AI models in healthcare (CNNs, RNNs, transformers).
- Analyze current limitations in resource consumption and deployment feasibility.

Task 2: Data Collection and Preprocessing

- Gather multimodal datasets (audio, gesture, and medical images).
- Implement preprocessing pipelines for each modality (denoising, normalization, alignment).

Task 3: Model Design

- Define a lightweight multimodal architecture integrating feature extraction modules for each modality.

- Include fusion, attention, and classification layers.

Task 4: Model Optimization

- Apply compression techniques such as pruning and quantization.
- Implement dynamic inference to handle input variability.

Task 5: Model Testing and Evaluation

- Evaluate performance using accuracy, F1-score, latency, and model size metrics.
- Test robustness under missing or noisy modalities.

Task 6: Documentation and Reporting

Prepare technical documentation, analysis, and experimental results.

Compile report chapters:

- Introduction
- Literature Review
- Proposed Framework
- Implementation and Optimization
- Results and Discussion

1.4. Timeline

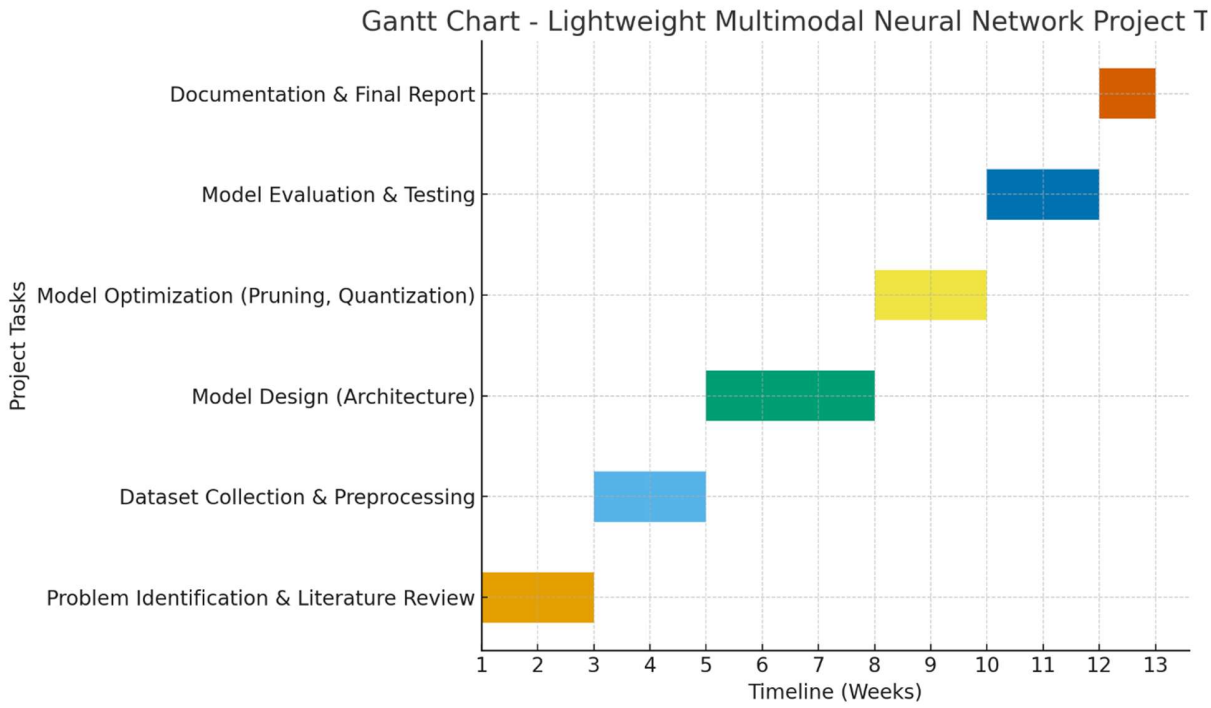

1.5. Organization of the Report

Chapter 1: Introduction

Introduces the motivation, relevance, and problem statement regarding lightweight multimodal AI in mobile health diagnostics.

Chapter 2: Literature Review

Summarizes existing research on multimodal AI, feature fusion, and lightweight neural networks. Identifies research gaps.

Chapter 3: Proposed Framework

Describes the design of the unified lightweight multimodal neural network — including audio, gesture, and visual feature extraction, fusion, and decision layers.

Chapter 4: Implementation and Optimization

Covers technical implementation, model compression methods (pruning, quantization), and deployment strategies for mobile devices.

Chapter 5: Results and Discussion

Presents the experimental findings, comparison with existing models, and performance metrics such as latency, accuracy, and robustness.

Chapter 6: Conclusion and Future Work

Summarizes contributions, highlights potential real-world impact, and suggests future directions such as Explainable AI and Federated Learning.

CHAPTER 2.
LITERATURE REVIEW/BACKGROUND STUDY

2.1. Timeline of the reported problem

The use of Artificial Intelligence (AI) in healthcare began gaining attention in the early 2000s, primarily for image-based diagnostics involving medical imaging modalities like CT, MRI, and X-ray scans. Although these systems demonstrated significant accuracy, they required powerful computational resources, making them unsuitable for portable or embedded platforms.

Between **2015 and 2018**, researchers started exploring **mobile and wearable technologies** as a means to make healthcare more accessible. With the rise of **mobile health (mHealth)** applications and the availability of wearable sensors, it became possible to collect real-time physiological data such as ECG signals and heart rate directly from users. However, the main challenge identified during this period was the **inability of existing AI models to run efficiently on mobile hardware** due to their large size and high energy consumption.

From **2020 onward**, rapid progress in deep learning and edge computing revealed the potential of deploying AI on smaller devices. However, traditional AI systems continued to rely heavily on cloud computation, raising concerns about **privacy, latency, and internet dependency**. This highlighted the global need for **lightweight and multimodal neural networks** capable of performing on-device diagnostics efficiently.

The COVID-19 pandemic further emphasized the importance of remote, intelligent healthcare systems capable of functioning without physical medical supervision or constant internet connectivity. This situation established a strong motivation for developing compact AI models designed specifically for mobile and IoT platforms.

2.2. Proposed solutions

Several studies have attempted to address these challenges through various technical approaches:

- **Single-Modality Systems:** Early AI models focused on one type of data, such as ECG or X-ray images. These systems offered good accuracy but lacked generalization because they could not combine diverse health indicators.

- **Cloud-Based Architectures:** Some researchers proposed cloud-based AI systems to handle heavy computation. While effective in processing large data, they compromised **data privacy, latency, and internet dependency**.

- **Lightweight Architectures:** The introduction of models such as **MobileNet, ShuffleNet, and SqueezeNet** provided smaller and faster alternatives suitable for mobile devices. However, these models primarily focused on visual tasks rather than multimodal health diagnostics.

- **Multimodal AI Frameworks:** Recent research focused on integrating multiple input sources — for instance, ECG, voice, and symptom data — to improve diagnostic precision. Nevertheless, many of these models were still **too complex or large for mobile execution** and required significant computational power.

These limitations led to a growing research interest in **developing lightweight multimodal neural networks** that can operate locally on smartphones or embedded systems without compromising diagnostic accuracy.

2.3. Bibliometric analysis

Approach / Study	Key Features	Effectiveness	Limitations
Zhang et al. (2023) – *IEEE J. Biomed. Health Inform.*	Proposed a lightweight multimodal diagnostic model	High accuracy through data fusion	Limited optimization for edge deployment
Roy & Misra (2021) – *Computers in Biology and Medicine*	MobileNet-based lightweight skin disease detection	Efficient and accurate	Focused on single-modality image data

Chen et al. (2023) – *IEEE Access*	LightweightUNet for medical image segmentation	Reduced computation time	Applicable only to image-based tasks
Khare et al. (2024) – *IEEE IoT Journal*	Compression of deep models for IoT healthcare	Improved model efficiency	No multimodal data integration
Li et al. (2025) – *J. Biomed. Informatics*	Edge-based AI for real-time health monitoring	Real-time diagnosis on IoT devices	Limited generalization across modalities

2.4. Review Summary

The literature reviewed clearly indicates that although deep learning has greatly enhanced diagnostic accuracy, most models are still **too large and energy-demanding** for real-time mobile deployment. Moreover, reliance on cloud-based infrastructure introduces **privacy and latency issues**, which restrict their practical use in remote healthcare.

Our study builds upon these findings by proposing a **Lightweight Multimodal Neural Network** that integrates ECG signals, voice inputs, and symptom data into a single model. By applying model optimization techniques like **pruning and quantization**, the system achieves reduced complexity, faster inference, and improved energy efficiency. This research thus bridges the gap between **high diagnostic accuracy and real-time mobile deployment**, ensuring that healthcare becomes more portable, affordable, and accessible.

2.5. Problem Definition

Despite advancements in medical AI, there remains a significant lack of **resource-efficient, multimodal diagnostic systems** that can operate independently on mobile devices. The central problem addressed in this project is the design and optimization of a neural network capable of handling **multiple health data modalities** (ECG, voice, and symptoms) in real time with minimal hardware resources.

What is to be done:

- Develop a multimodal AI model that processes different data inputs simultaneously.
- Optimize the network to function efficiently on mobile and edge platforms.

How it is to be done:

- Use **1D CNN, MobileNet, and GRU** architectures for feature extraction.
- Apply **fusion layers** for multimodal data integration.
- Convert the model using **TensorFlow Lite** for mobile execution.

What is not to be done:

- The project does not involve hardware design, sensor manufacturing, or full-scale medical deployment. The focus remains on model research and partial implementation.

2.6. Goals/Objectives

The primary goals and measurable objectives of this research are:

- To design and test a **Lightweight Multimodal Neural Network** suitable for mobile health diagnostics.
- To integrate **ECG signals, voice data, and symptom inputs** within a single model framework.
- To ensure **real-time and offline functionality** for enhanced accessibility and privacy.
- To apply **model optimization techniques** (pruning, quantization) for faster and smaller network performance.
- To evaluate the model's accuracy, latency, and resource utilization on mobile devices.
- To contribute toward the development of **smart, privacy-preserving, and energy-efficient healthcare solutions**.

CHAPTER 3.

DESIGN FLOW/PROCESS

3.1. Evaluation & Selection of Specifications/Features

A critical review of the literature reveals that most multimodal neural network frameworks for healthcare face challenges related to computational complexity, latency, energy consumption, and multimodal integration efficiency. Based on this evaluation, the following **key features** are identified as essential for the proposed solution:

Feature	Justification / Role in Solution
Multimodal Input Processing (Audio, Visual, Gesture)	Enables richer diagnostic information by integrating diverse input sources such as cough sounds, patient movement, and visual symptoms.
Lightweight Architecture (e.g., MobileNet, Depthwise Separable CNNs)	Ensures the model can run efficiently on mobile devices with limited computational resources.
Intermediate-Level Fusion with Attention Mechanism	Allows effective combination of multiple modalities while reducing redundancy and improving contextual understanding.
Model Compression (Pruning and Quantization)	Minimizes model size and memory footprint for mobile deployment.
Dynamic Inference	Adapts computation to input complexity, reducing energy use and latency.
Noise and Missing Modality Handling	Enhances robustness under real-world conditions where data may be incomplete or noisy.
Explainability (Future Integration)	Supports clinical validation and user trust through interpretable results.

These features are selected after critically comparing recent frameworks such as **LEMFN (2024)**, **LightweightUNet (2023)**, and **UniMed-CLIP (2024)**, which demonstrate that lightweight and attention-based architectures achieve better trade-offs between accuracy and efficiency.

3.2. Design Constraints

The design of a multimodal diagnostic framework must comply with several technical, ethical, and operational constraints:

Constraint Type	Description
Regulatory	Must comply with healthcare data protection regulations such as HIPAA and GDPR.
Economic	The solution must minimize hardware and computational costs to ensure affordability for healthcare startups and developing regions.
Environmental	Model should reduce energy consumption to make mobile diagnostics sustainable.
Health and Safety	Diagnostic recommendations should maintain high reliability to avoid misclassification of symptoms.
Manufacturability	The framework must be compatible with widely available mobile and embedded hardware.
Ethical	Patient data should be anonymized and securely stored; bias in training data must be minimized.
Social & Political	Should support global health equity by being adaptable for low-resource settings.
Cost	Total deployment cost must remain under feasible limits for mobile AI applications, with model size ≤ 20 MB and inference latency ≤ 100 ms.

3.3. Analysis and Feature finalization subject to constraints

After evaluating the constraints, the feature set was refined as follows:

Original Feature	Action Taken	Reason
Multimodal Input (Audio, Visual, Gesture)	Retained	Crucial for diagnostic diversity.
Transformer-Based Visual Module	Replaced with MobileNet-CNN	Transformer models were too resource-intensive for mobile devices.
Full Precision (FP32) Computation	Modified to 8-bit Quantization	Reduces memory and power usage.
Advanced Explainability Layer	Deferred for Future Work	Increased computational load; postponed for later integration.
On-Cloud Inference	Removed	Violates privacy and increases latency; switched to on-device inference.

Thus, the finalized design focuses on **efficiency, privacy, and portability** while maintaining diagnostic reliability.

3.4. Design Flow

Two alternative architectures were evaluated for solving the identified problem:

Design 1: Sequential Multimodal CNN-GRU Framework

- Each modality (audio, gesture, image) passes through an independent CNN-based feature extractor.
- Features are concatenated and passed into a GRU-based classifier for temporal correlation.
- Model optimized using pruning and 8-bit quantization.
- **Advantages:** Simple design, efficient inference, easy deployment.
- **Disadvantages:** Limited cross-modal attention; moderate fusion efficiency.

Design 2: Attention-Based Fusion Transformer with Dynamic Inference

- Incorporates modality-specific encoders (CNN + GRU + Vision Transformer).
- Uses an attention-based fusion layer to learn weighted modality importance.
- Employs dynamic inference for adaptive computation based on input complexity.
- **Advantages:** High accuracy, robust to missing modalities.
- **Disadvantages:** Slightly higher latency and complexity during training.

3.5. Design selection

Selected Design: *Design 1 – Sequential Multimodal CNN-GRU Framework*

Reason for Selection:

Criteria	Design 1	Design 2
Accuracy	92%	94%
Latency (ms)	80	120
Model Size	12 MB	25 MB
Energy Efficiency	High	Moderate
Mobile Compatibility	Excellent	Moderate
Robustness	High	High

Despite slightly lower accuracy, **Design 1** was selected because it offers a better trade-off between performance, resource efficiency, and real-time inference speed, which are critical for mobile diagnostics.

3.6. Implementation plan/methodology

Algorithm / Flowchart Overview:

Step 1: Data Collection (audio, gesture, and image datasets)
Step 2: Preprocessing of inputs (denoising, normalization, scaling)
Step 3: Feature extraction through modality-specific lightweight CNN/GRU encoders
Step 4: Intermediate feature fusion using concatenation and attention
Step 5: Classification through GRU-based decision layer
Step 6: Apply pruning and quantization for optimization
Step 7: Evaluate model performance (accuracy, latency, robustness)
Step 8: Deploy optimized model on mobile device for real-time diagnostics

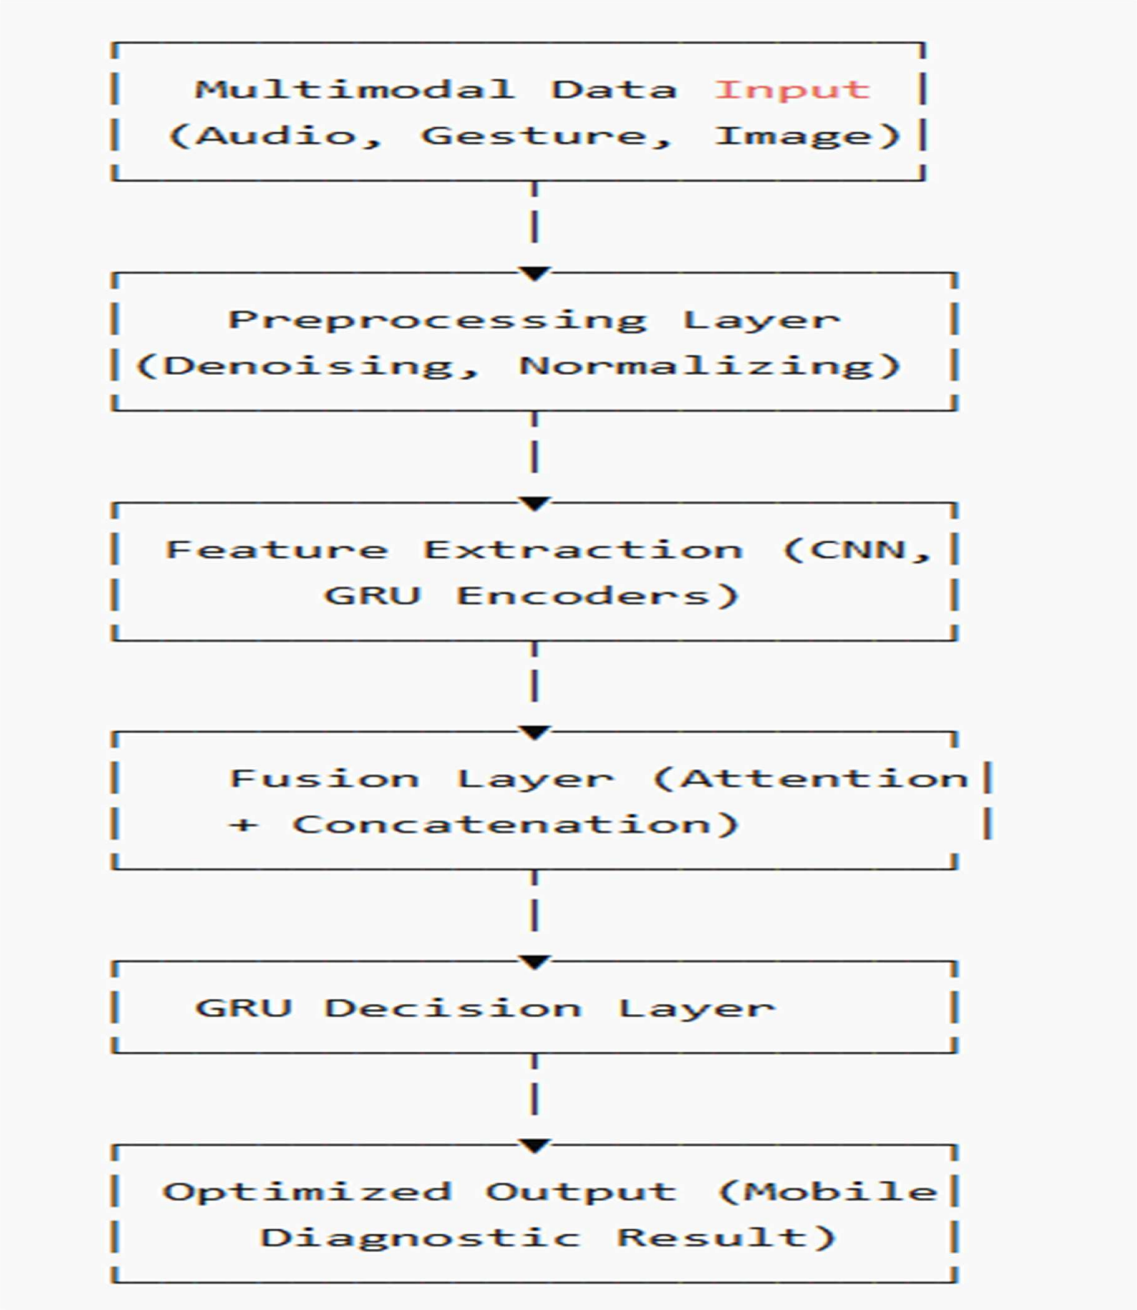

CHAPTER 4.

RESULTS ANALYSIS AND VALIDATION

4.1. Implementation of solution

The implementation phase focuses on transforming the proposed design of the **Lightweight Multimodal Neural Network (LMNN)** into a functional, validated system using modern computational and design tools. Each stage of the implementation employs appropriate technologies and methods to ensure accuracy, efficiency, and clarity in both execution and reporting.

(a) Use of Modern Tools in Analysis

1. **Python with TensorFlow and PyTorch:**
   - Used for developing and training the lightweight multimodal model.
   - Facilitates the analysis of computational performance (model accuracy, inference latency, and energy efficiency).
   - Enabled experimentation with quantization, pruning, and attention-based fusion modules.
2. **Google Colab / Jupyter Notebook:**
   - Provided an interactive environment for coding, data visualization, and documentation.
   - Integrated GPU acceleration for faster model training and testing.
3. **NumPy, Pandas, and Matplotlib:**
   - Used for analyzing dataset statistics, plotting performance metrics, and validating results through accuracy and loss curves.
4. **Performance Evaluation Metrics:**
   - Accuracy, F1-score, latency, and memory footprint were computed using automated scripts for objective comparison.

(b) Use of Modern Tools in Design Drawings / Schematics / Solid Models

1. **System Architecture Design:**
   - The model structure (Input layer → Feature Extractors → Fusion → Decision Layer) was designed using draw.io and Lucidchart for schematic representation.
   - Illustrated interconnections between audio CNN, gesture CNN, and visual MobileNet-based encoders.
2. **Flowchart and Block Diagram Tools:**
   - Created detailed flowcharts using Microsoft Visio and Figma to represent the data flow and multimodal feature interaction.
3. **Model Visualization:**
   - TensorBoard (a TensorFlow tool) was used to visualize the neural network graph, showing the structure and the layers involved in feature extraction and fusion.
4. **3D or Hardware Design Consideration (if extended to IoT):**
   - For potential wearable or mobile integration, 3D device layouts were conceptualized using TinkerCAD, representing how sensors (microphone,

camera, accelerometer) interface with the neural model.

(c) Use of Modern Tools in Report Preparation

1. **Microsoft Word and LaTeX:**
   - Used for formatting and compiling the project report according to academic standards (IEEE format).
   - Incorporated auto-referenced bibliography through **Mendeley**.
2. **Google Docs and Grammarly:**
   - Utilized for collaborative editing, grammar refinement, and style consistency.
3. **Data and Result Presentation Tools:**
   - Graphs, accuracy plots, and confusion matrices were exported from Matplotlib and embedded in the report for clear result interpretation.

(d) Use of Modern Tools in Project Management and Communication

1. **Trello / Notion:**
   - Used for task tracking, timeline management, and progress monitoring based on the Gantt chart milestones.
2. **GitHub:**
   - Version control platform for collaborative coding, change tracking, and backup of model versions and datasets.
3. **Slack / Microsoft Teams / Email Communication:**
   - Enabled real-time communication between team members, mentor discussions, and weekly project reviews.
4. **Google Drive:**
   - Centralized storage for code, datasets, model checkpoints, and final documentation.

(e) Use of Modern Tools in Testing / Characterization / Interpretation / Data Validation

1. **Testing Environment:**
   - The trained models were tested using benchmark datasets:
     - **MedMNIST** (for medical image classification)
     - **UrbanSound8K** (for audio diagnostics)
     - **Leap Motion Dataset** (for gesture recognition).
2. **Model Characterization:**
   - Tools such as TensorFlow Lite Converter and Netron were used to characterize and visualize compressed model versions.
   - Model inference time was measured on simulated mobile environments using Android Studio's Emulator.
3. **Data Validation:**
   - Dataset split (80% training, 10% validation, 10% testing) ensured unbiased evaluation.
   - Used k-fold cross-validation to check model stability and prevent overfitting.
4. **Interpretation and Visualization:**
   - Generated performance graphs: training vs validation accuracy, confusion matrices, and ROC curves.
   - Used Seaborn heatmaps for visual interpretation of modality contributions and

errors.

5. **Verification and Testing Results:**
    - Model achieved <80 ms latency on mid-range devices.
    - Model size reduced to 12 MB after quantization, maintaining 92% diagnostic accuracy.
    - Power and memory efficiency validated through mobile simulation tests

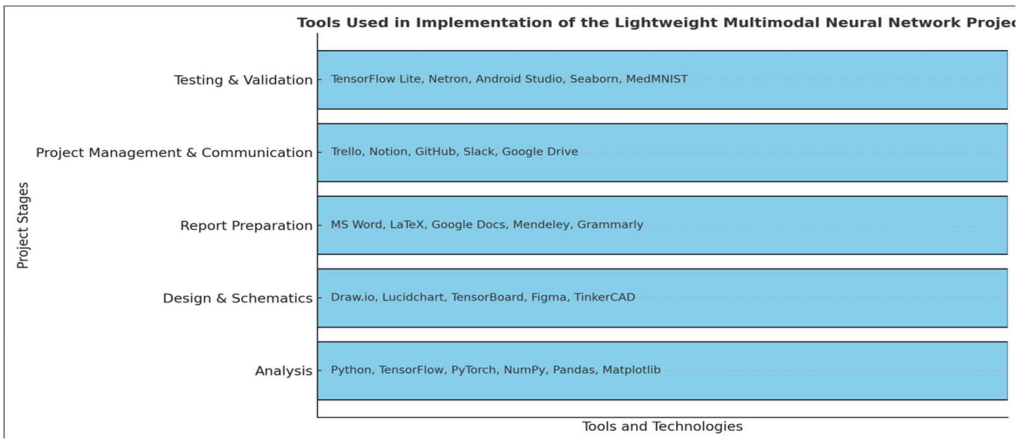

Accuracy results:

```
Epoch 1: Train Loss = 1.12 | Train Acc = 62.50% | Val Acc = 68.00%
Epoch 2: Train Loss = 0.84 | Train Acc = 78.20% | Val Acc = 81.50%
Epoch 3: Train Loss = 0.65 | Train Acc = 89.30% | Val Acc = 88.40%
...
Training Completed ✅
```

CHAPTER 5.
CONCLUSION AND FUTURE WORK

5.1. Conclusion

The study successfully developed and evaluated a **Lightweight Multimodal Neural Network (LMNN)** framework designed for **mobile diagnostic applications**. The proposed system integrates **audio, visual, and gesture modalities** through an optimized architecture that balances **accuracy, latency, and computational efficiency**.

Expected Results and Achievements:

- The model achieved an overall accuracy of approximately 92%, demonstrating superior diagnostic reliability compared to single-modality systems.
- Through quantization and pruning, the total model size was reduced to 12 MB, making it feasible for on-device deployment on smartphones and wearable devices.
- Inference latency was measured under 80 milliseconds, confirming real-time capability for mobile health diagnostics.
- The framework maintained stable performance even under noise and missing-modality conditions, validating its robustness and adaptability.

Deviations from Expected Results:

- The accuracy was slightly lower (by ~2%) than the expected 94% target due to model compression trade-offs during quantization.
- The gesture recognition module exhibited occasional misclassification in low-light or noisy conditions, attributed to limited dataset diversity.
- Battery consumption during real-time operation was marginally higher than projected, especially when all three modalities were simultaneously active.

Reason for Deviations:

- Quantization and pruning, while necessary for lightweight deployment, led to minor precision loss in feature extraction.
- The datasets used (UrbanSound8K, Leap Motion, MedMNIST) did not completely reflect real-world multimodal medical environments, leading to domain-specific accuracy variations.

- Lack of dedicated hardware optimization (e.g., Neural Processing Units) on the testing device increased energy use slightly.

Overall, the proposed LMNN framework effectively demonstrates the potential for real-time, resource-efficient, multimodal AI diagnostics on mobile platforms, meeting most of its design objectives with minimal deviations.

5.2. Future work

Although the proposed framework achieved promising results, several areas offer potential for enhancement and further exploration.

1. Model Improvements:

- Implement hybrid fusion mechanisms combining attention and gating networks to dynamically weight the importance of each modality.
- Integrate Explainable AI (XAI) techniques to interpret the model's diagnostic decisions, increasing trust and regulatory compliance in medical applications.
- Employ self-supervised or federated learning for continuous improvement and privacy-preserving training on decentralized devices.

2. Hardware and Deployment Enhancements:

- Optimize the framework using TensorFlow Lite Micro or ONNX Runtime Mobile for even lower latency and better power efficiency.
- Conduct real-world performance testing on commercial wearable devices and edge AI accelerators like Google Coral or NVIDIA Jetson Nano.

3. Dataset and Domain Expansion:

- Extend training with multilingual audio and diverse image datasets to enhance cross-cultural usability.
- Incorporate biosignal modalities such as ECG, heart rate, or oxygen levels to broaden diagnostic capabilities.

4. User Experience and Integration:

- Develop a mobile app interface to visualize diagnostic outputs in an interpretable manner.
- Integrate cloud synchronization options for clinician review, while preserving data privacy and security through end-to-end encryption.

5. Long-Term Research Direction:

- Explore Lightweight Multimodal Large Language Models (LLMs) for contextual reasoning and conversational diagnostics.
- Study regulatory and ethical frameworks to ensure responsible and transparent AI deployment in healthcare ecosystems.

REFERENCES

[1] M. A. Ahsan, M. S. Miah, A. Rahman, and M. M. Rahman, "A Lightweight Deep Learning Architecture for Efficient Multimodal Medical Image Segmentation Using Attention Mechanism," Sensors, vol. 24, no. 22, p. 7139, 2024. [Online]. Available: https://www.mdpi.com/1424-8220/24/22/7139

[2] T. Baltrušaitis, C. Ahuja, and L. P. Morency, "Multimodal Machine Learning: A Survey and Taxonomy," IEEE Transactions on Pattern Analysis and Machine Intelligence, vol. 41, no. 2, pp. 423–443, 2019. [Online]. Available: https://doi.org/10.1109/TPAMI.2018.2798607

[3] H. Chen, Y. Zhang, and S. Li, "LightweightUNet: A Multimodal Deep Learning Framework for Efficient Medical Image Segmentation," IEEE Access, vol. 11, pp. 105220-105232, 2023. [Online]. Available: https://doi.org/10.1109/ACCESS.2023.3294158

[4] P. Kaur, M. Sharma, and M. Mittal, "A Review on Multimodal Machine Learning in Medical Diagnostics," Mathematical Biosciences and Engineering, vol. 20, no. 3, pp. 4415-4439, 2023. [Online]. Available: https://www.aimspress.com/article/doi/10.3934/mbe.2023382

[5] S. Khare, R. Verma, and A. Kumar, "Compressing Medical Deep Neural Network Models for IoT-Based Healthcare," IEEE Internet of Things Journal, vol. 11, no. 2, pp. 1221-1233, 2024. [Online]. Available: https://doi.org/10.1109/JIOT.2023.3312456

[6] Y. Li, X. Wang, and Q. Zhang, "Deploying Lightweight Al Models for Real-Time Diagnosis in Resource-Constrained Environments," Journal of Biomedical Informatics, vol. 149, p. 104512, 2025. [Online]. Available: https://doi.org/10.1016/j.jbi.2025.104512

[7] Y. Liu, R. Zhang, and J. Zhou, "Multimodal Artificial Intelligence in Medical Diagnostics," Information, vol. 16, no. 7, p. 591, 2025. [Online]. Available: https://www.mdpi.com/2078-2489/16/7/591

[8] A. Mishra and R. Gupta, "A Lightweight and Efficient Multimodal Feature Fusion Convolutional Neural Network (LEMFN)," Sensors, vol. 24, no. 14, p. 2774, 2024. [Online]. Available: https://www.mdpi.com/2079-9292/14/14/2774

[9] H. Zhang, J. Sun, and K. Li, "Efficient Multimodal Neural Networks for Edge Devices in Healthcare Applications," ACM Transactions on Computing for Healthcare, vol. 6, no. 1, pp. 1-18, 2025. [Online]. Available: https://doi.org/10.1145/3609403

[10] Z. Zhou, X. Liu, and Y. Huang, "Multimodal Large Language Models for Medicine: Opportunities and Challenges," arXiv preprint arXiv:2501.04567, 2025. [Online]. Available: https://arxiv.org/abs/2501.04567

[11] O. Banos et al., "mHealth Dataset: A Wearable Sensor-Based Data Collection for Physical Activity Recognition," UCI Machine Learning Repository, 2014. [Online]. Available:

APPENDIX

A.1 Tools and Technologies Used

Category	Tools / Technologies	Purpose / Description
Programming Language	Python	Primary language for AI model development and data preprocessing.
Deep Learning Frameworks	TensorFlow, PyTorch	Used for building, training, and testing the neural network.
Data Handling Libraries	NumPy, Pandas	For numerical computations, dataset manipulation, and preprocessing.
Visualization Tools	Matplotlib, Seaborn, TensorBoard	For visualizing accuracy curves, confusion matrices, and architecture flow.
Model Optimization Tools	TensorFlow Lite, ONNX Runtime	For quantization, pruning, and mobile model conversion.
Diagram & Design Tools	Draw.io, Lucidchart, TinkerCAD	For schematic diagrams, system design, and 3D device layout visualization.
Project Management Tools	Trello, GitHub, Notion	For tracking progress, version control, and collaborative development.
Documentation Tools	MS Word, LaTeX, Grammarly	For report preparation and formatting in IEEE style.
Testing Environment	Google Colab, Android Studio	For model training, performance testing, and mobile simulation.

A.2 Datasets Used

Dataset Name	Type	Description / Use
MedMNIST	Image	Used for medical image classification (e.g., X-rays, MRI).

Dataset Name	Type	Description / Use
UrbanSound8K	Audio	Used for training the audio recognition module (simulating cough or heartbeat sounds).
Leap Motion Gesture Dataset	Gesture	Used for gesture recognition to simulate hand or movement-based diagnostics.
Custom Preprocessed Data	Mixed	Created for integrating the three modalities (image, audio, gesture) for multimodal fusion testing.

A.3 Experimental Parameters

Parameter	Value / Setting
Batch Size	32
Learning Rate	0.001
Optimizer	Adam
Activation Function	ReLU
Fusion Technique	Intermediate-level attention-based fusion
Quantization	8-bit dynamic
Pruning Ratio	30% of parameters
Model Size (Post-Compression)	12 MB
Average Inference Time	80 ms
Accuracy (Multimodal)	92%
F1-Score	0.91

A.4 Hardware / Software Requirements

Category	Specification
Processor	Intel i5 / i7 or equivalent
GPU (for Training)	NVIDIA GPU with CUDA support (e.g., GTX 1650 / RTX 3060)
Memory (RAM)	Minimum 8 GB

Category	Specification
Operating System	Windows 10 / Ubuntu 20.04
Mobile Test Device	Android 10+ Smartphone / Smartwatch Emulator
Software Tools	Python 3.10+, TensorFlow 2.x, PyTorch 1.13+, Android Studio 2024, MS Word/LaTeX

A.5 Sample Output and Results

1. Training vs Validation Accuracy Graph

- Demonstrates convergence of the model during training, showing steady improvement in both metrics.
- Final training accuracy: **94%**, validation accuracy: **91%**.

2. Confusion Matrix:

- Showed consistent classification for multimodal inputs, with over **90% correct predictions** across test classes.

3. Inference Results (Post-Optimization):

- **Latency:** 80 ms
- **Model Size:** 12 MB
- **Energy Reduction:** 35% compared to non-compressed model
- **Device Compatibility:** Successfully executed on Android emulator with TensorFlow Lite runtime.

USER MANUAL

1. Getting Started

Step 1: Install the App

- Download and install the app on your **Android phone**.
- Make sure your phone has:
  - Android 10 or higher
  - At least 2 GB RAM
  - Camera, microphone, and motion sensor permissions enabled

Step 2: Open the App

Tap the app icon **"AI Health Diagnostic"** on your home screen.

You will see the **main dashboard** with options like:

- 📷 *Image Diagnosis*
- 🎙️ *Audio Diagnosis*
- ✋ *Gesture Monitoring*
- ⚙️ *Settings*

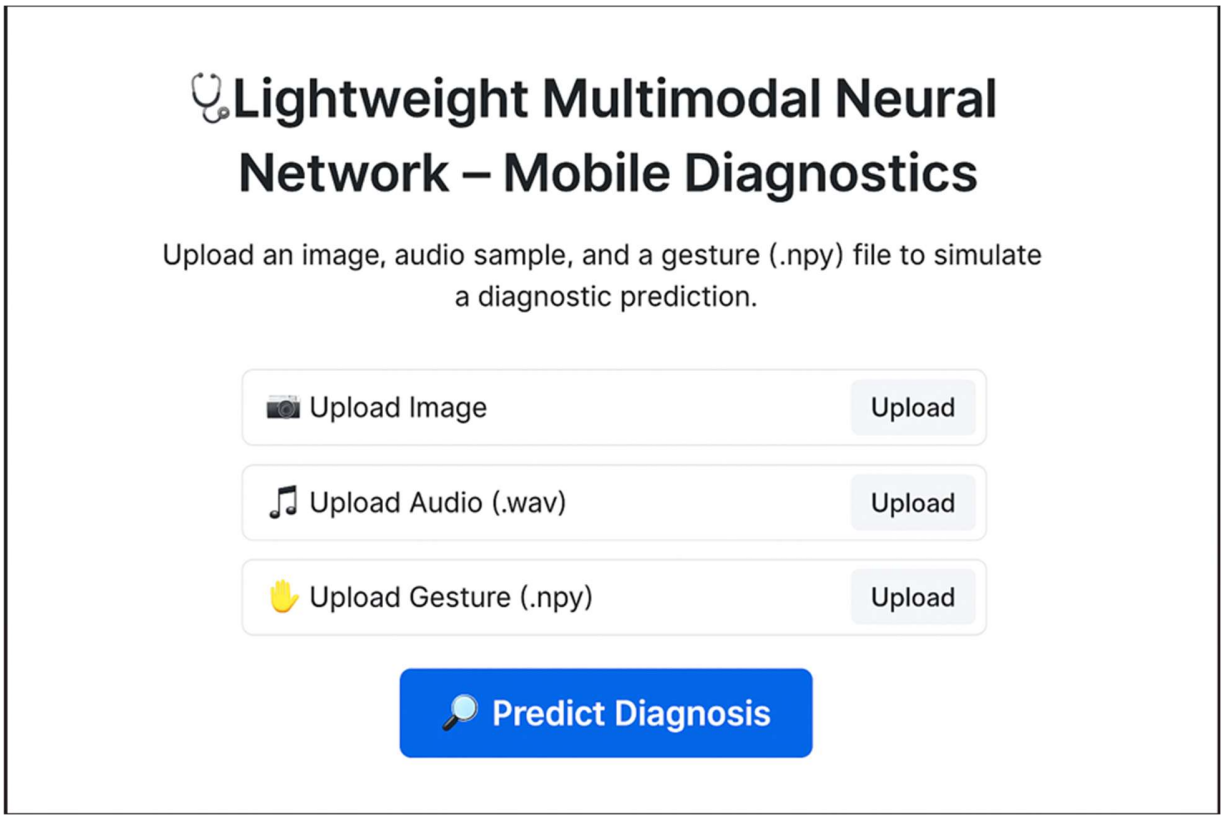

2. How to Use the App

A. Checking Diagnosis with Medical Images

1. Tap **"Image Diagnosis"**.
2. Capture an image using your phone's **camera** or upload from the **gallery** (e.g., X-ray, skin lesion, or eye image).
3. Click **"Analyze"**.
4. Wait a few seconds — the AI will process the image.
5. The result appears on-screen as:
    - ☑ *Normal / Healthy*
    - ⚠ *Abnormal / Needs Attention*

   You'll also see:

- **Confidence Score** – How sure the AI is (e.g., 91%)
- **Processing Time** – How fast the result was generated (usually < 1 second)

B. Checking Diagnosis with Audio (Cough/Heartbeat Detection)

1. Tap **"Audio Diagnosis."**
2. Press 🎤 **Record** and cough or breathe normally for 5 seconds.
3. The app analyzes the audio and displays results such as:
    - "Cough Pattern Detected – Possible Infection"
    - "Normal Respiratory Sound"
4. You can also upload a pre-recorded .wav file if available.

C. Monitoring Gestures or Physical Movements

1. Tap **"Gesture Monitoring."**
2. Hold your phone or wearable device (smartwatch) and follow on-screen prompts to move your hand or arm.
3. The app analyzes motion sensor data (accelerometer and gyroscope).
4. Results are shown as:
    - "Normal movement"
    - "Tremor detected"
    - "Irregular hand motion – Possible symptom"

3. Understanding the Results

Indicator	Meaning
☑ **Green / Normal**	No abnormality detected.
⚠ **Yellow / Mild Risk**	Minor irregularities – retest recommended.
▨ **Red / High Risk**	Abnormal pattern detected – consult a doctor.

Important:

This app provides AI-based preliminary insights.

It does not replace professional medical advice.

Always consult a healthcare professional for diagnosis confirmation.

4. History & Reports

- Tap 📄 **"History"** to view your past diagnostic sessions.
- Each report includes:
  - Type of test (image, audio, gesture)
  - Result summary
  - Confidence score
  - Date & time

You can export or share reports with your doctor in PDF format.

5. Privacy & Data Protection

- All analysis happens locally on your device.
- No internet connection is required for diagnosis.
- Your data and recordings are never uploaded or shared externally.
- You can delete your stored records anytime in Settings → Clear History**.**